

PROBLEM STATEMENT:- TO PREDICT THE RAINFALL BASED ON VARIOUS FEATURES OF THE DATASET

Linear Regression

Data Collection:

```
In [1]: import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv(r"C:\Users\91903\Downloads\rainfall in india 1901-2015.csv")
df
```

Out[2]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4

4116 rows × 19 columns



Data Cleaning and preprocessing:

In [3]:

df.head()

Out[3]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4

In [4]:

df.tail()

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	1
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	2

In [5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBDIVISION     4116 non-null   object
1   YEAR            4116 non-null   int64
2   JAN             4112 non-null   float64
3   FEB             4113 non-null   float64
4   MAR             4110 non-null   float64
5   APR             4112 non-null   float64
6   MAY             4113 non-null   float64
7   JUN             4111 non-null   float64
8   JUL             4109 non-null   float64
9   AUG             4112 non-null   float64
10  SEP             4110 non-null   float64
11  OCT             4109 non-null   float64
12  NOV             4105 non-null   float64
13  DEC             4106 non-null   float64
14  ANNUAL          4090 non-null   float64
15  Jan-Feb         4110 non-null   float64
16  Mar-May         4107 non-null   float64
17  Jun-Sep         4106 non-null   float64
18  Oct-Dec         4103 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [6]: df.describe()

Out[6]:

	YEAR	JAN	FEB	MAR	APR	MAY	JUN
count	4116.000000	4112.000000	4113.000000	4110.000000	4112.000000	4113.000000	4111.000000
mean	1958.218659	18.957320	21.805325	27.359197	43.127432	85.745417	230.234444
std	33.140898	33.585371	35.909488	46.959424	67.831168	123.234904	234.710755
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.400000
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.350000
50%	1958.000000	6.000000	6.700000	7.800000	15.700000	36.600000	138.700000
75%	1987.000000	22.200000	26.800000	31.300000	49.950000	97.200000	305.150000
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.900000

```
In [8]: df.isna().any()
```

```
Out[8]: SUBDIVISION    False
        YEAR           False
        JAN            True
        FEB            True
        MAR            True
        APR            True
        MAY            True
        JUN            True
        JUL            True
        AUG            True
        SEP            True
        OCT            True
        NOV            True
        DEC            True
        ANNUAL         True
        Jan-Feb        True
        Mar-May        True
        Jun-Sep        True
        Oct-Dec        True
        dtype: bool
```

```
In [9]: df.fillna(method='ffill',inplace=True)
```

```
In [10]: df.isnull().sum()
```

```
Out[10]: SUBDIVISION    0
        YEAR           0
        JAN            0
        FEB            0
        MAR            0
        APR            0
        MAY            0
        JUN            0
        JUL            0
        AUG            0
        SEP            0
        OCT            0
        NOV            0
        DEC            0
        ANNUAL         0
        Jan-Feb        0
        Mar-May        0
        Jun-Sep        0
        Oct-Dec        0
        dtype: int64
```

In [11]: `df.describe()`

Out[11]:

	Y	JUN	JUL	AUG	SEP	OCT	NOV	DEC	
count	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	41
mean	230.567979	347.177235	290.239796	197.524781	95.724198	40.081997	19.042225	14	
std	234.896056	269.321089	188.785639	135.509037	99.689878	68.851397	42.655830	9	
min	0.400000	0.000000	0.000000	0.100000	0.000000	0.000000	0.000000		
25%	70.475000	175.900000	155.850000	100.575000	14.600000	0.700000	0.100000	8	
50%	138.900000	284.800000	259.400000	174.000000	65.750000	9.700000	3.100000	11	
75%	306.150000	418.325000	377.800000	266.225000	148.600000	46.325000	17.600000	16	
max	1609.900000	2362.800000	1664.600000	1222.000000	948.300000	648.900000	617.500000	63	

In [12]: `df.columns`

Out[12]: Index(['SUBDIVISION', 'YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL',
'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb', 'Mar-May',
'Jun-Sep', 'Oct-Dec'],
dtype='object')

In [13]: `df.shape`

Out[13]: (4116, 19)

In [14]: `df['Jan-Feb'].value_counts()`

Out[14]: Jan-Feb
0.0 238
0.1 80
0.2 52
0.3 38
0.4 32
...
23.3 1
95.2 1
76.9 1
66.5 1
69.3 1
Name: count, Length: 1220, dtype: int64

```
In [15]: df['ANNUAL'].value_counts()
```

```
Out[15]: ANNUAL
790.5      4
770.3      4
1836.2     4
1024.6     4
1926.5     3
..
443.9      1
689.0      1
605.2      1
509.7      1
1642.9     1
Name: count, Length: 3712, dtype: int64
```

```
In [16]: df['Mar-May'].value_counts()
```

```
Out[16]: Mar-May
0.0        29
0.1        13
0.3        11
8.3        11
11.5       10
..
246.3      1
248.1      1
151.3      1
249.5      1
223.9      1
Name: count, Length: 2262, dtype: int64
```

```
In [17]: df['Jun-Sep'].value_counts()
```


```
Out[17]: Jun-Sep
434.3      4
334.8      4
573.8      4
613.3      4
1082.3     3
..
301.6      1
380.9      1
409.3      1
229.4      1
958.5      1
Name: count, Length: 3683, dtype: int64
```

```
In [18]: df['Oct-Dec'].value_counts()
```

```
Out[18]: Oct-Dec
0.0      16
0.1      15
0.5      13
0.6      12
0.7      11
..
191.5     1
124.5     1
139.1     1
41.5      1
555.4     1
Name: count, Length: 2389, dtype: int64
```

Exploratory Data analysis

```
In [19]: df=df[['JAN', 'FEB', 'MAR', 'APR', 'DEC']]
sns.heatmap(df.corr(),annot=True)
plt.show()
```



```
In [20]: df.columns
```

```
Out[20]: Index(['JAN', 'FEB', 'MAR', 'APR', 'DEC'], dtype='object')
```



```
In [21]: x=df[["FEB"]]  
y=df["JAN"]
```

```
In [22]: from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state
```

```
In [23]: from sklearn.linear_model import LinearRegression  
reg=LinearRegression()  
reg.fit(X_train,y_train)  
print(reg.intercept_)  
coeff_=pd.DataFrame(reg.coef_,x.columns,columns=['coefficient'])  
coeff_
```

9.650666612303553

```
Out[23]:
```

	coefficient
FEB	0.442278


```
In [24]: score=reg.score(X_test,y_test)  
print(score)
```

0.1793580786264921


```
In [25]: predictions=reg.predict(X_test)
```

```
In [26]: plt.scatter(y_test,predictions)
```

```
Out[26]: <matplotlib.collections.PathCollection at 0x2f43cf33d30>
```



```
In [27]: df500=df[:][:500]
sns.lmplot(x="FEB",y="JAN",order=2,ci=None,data=df500)
plt.show()
```




```
In [28]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
reg.fit(X_train,y_train)
reg.fit(X_test,y_test)
```

Out[28]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [29]: y_pred=reg.predict(X_test)
plt.scatter(X_test,y_test,color='blue')
plt.plot(X_test,y_pred,color='black')
plt.show()
```



```
In [30]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 Score:",r2)
```

R2 Score: 0.21084431570038997

Ridge Regression

```
In [31]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [32]: features= df.columns[0:5]
target= df.columns[-5]
```

```
In [33]: x=np.array(df['JAN']).reshape(-1,1)
y=np.array(df['FEB']).reshape(-1,2)
```

```
In [34]: x= df[features].values
y= df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=
```

```
In [35]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
```


```
In [36]: print("\n Ridge Model:\n")
print("the train score for ridge model is{}".format(train_score_ridge))
print("the test score for ridge model is{}".format(test_score_ridge))
```

Ridge Model:

the train score for ridge model is0.9999999999874192
the test score for ridge model is0.99999999998833

```
In [37]: lr=LinearRegression()
```

```
In [39]: plt.figure(figsize= (10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=10,color='red')
plt.plot(features,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



Lasso regression


```
In [40]: print("\n Lasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is{}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.9999207747038827
 The test score for ls model is 0.9999206791315255

```
In [41]: pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="bar")
```


Out[41]: <Axes: >



```
In [42]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],random_state=0).fit(x_train,y
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.9999999999999921
 0.9999999999999921

```
In [44]: plt.figure(figsize= (10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=6,color="red")
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color="blue")
plt.plot(features,alpha=0.4,linestyle='none',marker='o',markersize=7,color="green")
plt.xticks(rotation = 90)
plt.legend()
plt.show()
```



ElasticNet Regression

```
In [45]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
print(regr.score(x,y))
```

```
[9.99098574e-01 0.00000000e+00 3.02728910e-05 0.00000000e+00
 0.00000000e+00]
0.016258606966612632
0.9999992160905338
```

```
In [46]: y_pred_elastic = regr.predict(x_train)
mean_squared_error=np.mean((y_pred_elastic - y_train)**2)
print(mean_squared_error)
```

```
0.0008816302333951303
```

Conclusion:

For the given dataset, we have performed linear regression, ridge regression, lasso regression, elastic regression.

Linear Regression: -0.0001948420411366225

Ridge Regression: 0.9999999999897634

Lasso Regression: 0.9999999999999921

Elastic Net Regression: 0.9999992133148984

Among all the models we observed that Elastic Net Regression got highest accuracy.

```
In [ ]:
```